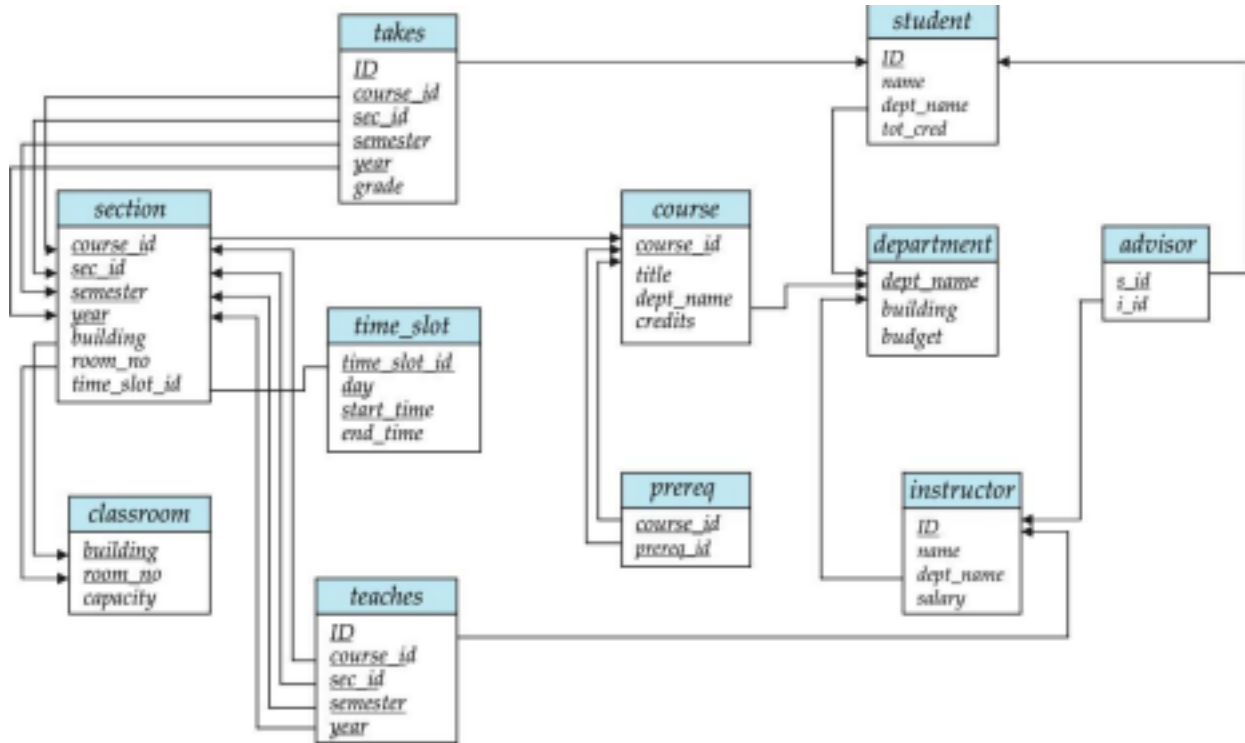


## WEEK - 2 Demo

### Queries on Single relations:

So far our example queries were on a single relation.

For Example: Consider a University Database



and the following query,

```
select name  
  
from instructor ;
```

where instructor table has attributes (ID, name, dept\_name, salary) Some important cases:

- If we want to force the elimination of duplicates, we insert the keyword **distinct** after **select**. We can rewrite the preceding query as:

```
select dept_name  
  
from instructor ;
```

- The **select** clause may also contain arithmetic expressions involving the operators +, -, \*, and / operating on constants or attributes of tuples. For example, the query,

```
select ID, name, dept_name, salary * 1.1  
  
from instructor ;
```

outputs a list of all the instructors, their department name and their salaries when increased by 10%.

### Queries on multiple relations:

Queries often need to access information from multiple relations. We now study how to write such queries.

1. Consider the University Database (ER Diagram) given above. Suppose we want to answer the query “Retrieve the names of all instructors, along with their department names and department building name.”

Looking at the schema of the relation *instructor*, we realize that we can get the department name from the attribute *dept\_name* but the department building name is present in the attribute *building* of the relation *department*.

To answer the above query, we list the relations that need to be accessed in the **from** clause, and specify the matching condition in the **where** clause.

```
select name, instructor.dept_name, building
from instructor, department
where instructor.dept_name = department.dept_name
```

(Note: The attribute *dept\_name* required the prefix as the relation’s name (instructor or department) and others did not. Why??)

2. To find instructor names and course identifiers for instructors in the Computer Science department, we could add an extra predicate to the **where** clause, as shown below.

```
select name, course_id
from instructor, teaches
where instructor.ID= teaches.ID and instructor.dept_name = 'Comp. Sci.';
```

### Aliasing or Renaming:

The basic syntax of table alias is as follows –

```
select column1, column2....
from table_name as alias_name
where [condition];
```

The basic syntax of column alias is as follows –

```
select column_name as alias_name
from table_name
where [condition];
```

**Example:** Consider the example 1 of the above section. The same query can be written as,

```
select name, dept_name as DepartmentName, building
from instructor as ins, department as dept
where ins.dept_name = dept.dept_name
```

## String Operations:

1. **String concatenation:** PostgreSQL allows you to directly concatenate strings, columns and int values using ' || ' operator. Here is the SQL query to concatenate columns first\_name and last\_name using ' || ' operator. You can even concatenate string with int using '||' operator

Example:

```
select first_name || ' ' || last_name as customer_name
from customer
limit 5;
```

2. **Convert string to lowercase:** the LOWER function is used to convert a string, an expression, or values in a column to lowercase.

**Syntax:** LOWER(string or value or expression)

Example:

```
select LOWER(title)
from film;
```

3. **Convert string to uppercase:** Like the LOWER function, the UPPER function accepts a string expression or string-convertible expression and converts it to an upper case format.

**Syntax:** UPPER(string\_expression)

Example:

```
select UPPER(title)
from film;
```

4. **Replace substring:** 'Replace' function searches and replaces a substring with a new substring in a string.

**Syntax:** **REPLACE**(source, old\_text, new\_text );

where;

- **source** is a string where you want to replace.
- **old\_text** is the text that you want to search and replace. If the *old\_text* appears multiple times in the string, all of its occurrences will be replaced.
- **new\_text** is the new text that will replace the old text (*old\_text*).

Example:

```
1. UPDATE table_name
   SET column_name = REPLACE(column,old_text,new_text)
   WHERE condition
```

```
2. select REPLACE('ABC AA', 'A', 'Z'); // output: ZBC ZZ
```

5. **Length of a string** is a function that returns the number of characters in a string. It is used to determine the length of a string in terms of the number of characters it contains.

**Syntax:** **LENGTH** (string\_expression)

Example:

```
select
LENGTH('hello');
this will return 5
```

6. **SPLIT\_PART** :- This function is used to split a string into parts based on a delimiter and extract a specific part.

Syntax - **SPLIT\_PART**(string, delimiter, part\_number)

Example - **SPLIT\_PART**('Big Adventure', ' ', 1)

This would return 'Big', because it's the first word before the space ( ' ' ).

**6. LPAD and RPAD :-** The LPAD function pads the left side of a string with a specified character until the string reaches a desired length. Similarly The RPAD function pads the right side of a string with a specified character until the string reaches a desired length.

Syntax - RPAD/LPAD(string, length, pad\_char)

Example - RPAD('Big', 10, '@') it will give - Big@@@@@

**6. Wildcard matching :-** % operator matches any sequence of characters, including an empty sequence.

Syntax -

SELECT column\_name FROM table\_name WHERE column\_name LIKE 'pattern';

Example -

words that **start with "J"**, such as John, Jane, Jack.

SELECT \* FROM employees WHERE name LIKE 'J%';

words that **end with "son"**, such as Johnson, Emerson.

SELECT \* FROM employees WHERE name LIKE '%son';

words that **contain "ar"**, such as Carl, Darren.

SELECT \* FROM employees WHERE name LIKE '%ar%';

## Ordering

**SELECT**

select\_list

**FROM**

table\_name

**ORDER BY**

sort\_expression1 [ASC | DESC],

...

sort\_expressionN [**ASC** | **DESC**];

When you query data from a table, the **SELECT** statement returns rows in an unspecified order. To sort the rows of the result set, you use the ORDER BY clause.

- First, specify a sort expression, which can be a column or an expression, that you want to sort after the ORDER BY keywords. If you want to sort the result set based on multiple columns or expressions, you need to place a comma (,) between two columns or expressions to separate them.
- Second, you use the ASC option to sort rows in ascending order and the DESC option to sort rows in descending order. If you omit the ASC or DESC option, the ORDER BY uses ASC by default.

Eg:- The below query

```
select ID, name, dept_name, salary * 1.1 as new_salary  
  
from instructor
```

```
Order by new_salary DESC, name ASC;
```

The result of the above query, the results will be first sorted by the new\_salary in the descending order, for those rows where new\_salary is the same, they will be sorted by the instructor name in ascending order.

## **Set Operations**

The SQL operations union, intersect, and except operate on relations and correspond to the mathematical set-theory operations  $\cup$ ,  $\cap$ , and  $-$ .

Eg:- The set of all courses taught in the Fall 2009 semester:

```
select course_id  
  
from section  
  
where semester = 'Fall' and year= 2009;
```

The set of all courses taught in the Spring 2010 semester:

```
select course_id
```

**from section**

**where semester = 'Spring' and year= 2010;**

In our discussion that follows, we shall refer to the relations obtained as the result of the preceding queries as c1 and c2, respectively.

<i>course_id</i>	<i>course_id</i>
CS-101	CS-101
CS-347	CS-315
PHY-101	CS-319
	CS-319
	FIN-201
	HIS-351
	MU-199

Output of c1, and c2 respectively.

## The Union Operation

To find the set of all courses taught either in Fall 2009 or in Spring 2010, or both, we write:

```
(select course_id
from section
where semester = 'Fall' and year= 2009)
union
(select course_id
from section
where semester = 'Spring' and year= 2010);
```

Output of above query:-

<i>course_id</i>
CS-101
CS-315
CS-319
CS-347
FIN-201
HIS-351
MU-199
PHY-101

The union operation automatically eliminates duplicates. If we want to retain all duplicates, we must write union all in place of union:

```
(select course_id
from section
where semester = 'Fall' and year= 2009)
union all
(select course_id
from section
where semester = 'Spring' and year= 2010);
```

The number of duplicate tuples in the result is equal to the total number of duplicates that appear in both c1 and c2. 3.5.2

## The Intersect Operation

To find the set of all courses taught in the Fall 2009 as well as in Spring 2010 we write:

```
(select course_id
from section
where semester = 'Fall' and year= 2009)
intersect
(select course_id
from section
where semester = 'Spring' and year= 2010);
```



Output of above query:-

<i>course_id</i>
CS-101

The intersect operation automatically eliminates duplicates. If we want to retain all duplicates, we must write intersect all in place of intersect:

```
(select course_id
from section
where semester = 'Fall' and year= 2009)
intersect all
(select course_id
from section
where semester = 'Spring' and year= 2010);
```

The number of duplicate tuples that appear in the result is equal to the minimum number of duplicates in both c1 and c2.

### **The Except Operation**

To find all courses taught in the Fall 2009 semester but not in the Spring 2010 semester, we write:

```
(select course_id
from section
where semester = 'Fall' and year= 2009)
except
(select course_id
from section
where semester = 'Spring' and year= 2010);
```

Output of above query:-

<i>course_id</i>
CS-347
PHY-101

The except operation outputs all tuples from its first input that do not occur in the second input; that is, it performs set difference. The operation automatically eliminates duplicates in the inputs before performing set difference. If we want to retain duplicates, we must write except all in place of except:

```
(select course_id
from section
where semester = 'Fall' and year= 2009)
except all
(select course_id
from section
where semester = 'Spring' and year= 2010);
```

The number of duplicate copies of a tuple in the result is equal to the number of duplicate copies in c1 minus the number of duplicate copies in c2, provided that the difference is positive.

### **The Outer Join :**

An outer join is a type of join operation in a relational database management system (RDBMS) that retrieves records from two or more tables and returns all matching rows from the joined tables, plus unmatched rows from one or both tables.

1. **Left Outer Join (or Left Join):** This type of join returns all rows from the left table (the first table mentioned in the query) and the matched rows from the right table. If there is no match found in the right table, NULL values are returned for the columns of the right table.
2. **Right Outer Join (or Right Join):** This type of join returns all rows from the right table and the matched rows from the left table. If there is no match found in the left table, NULL values are returned for the columns of the left table.

Syntax : -

***select \****

```
from table1  
  
left outer join table2  
  
on table1.id = table1.id;
```

outer join can be used as set difference or like except by using conditions filter with

it. ***select \****

```
from table1  
  
left outer join table2  
  
on table1.id = table1.id  
  
where table1.id is not NULL;
```

this query will filter out rows where id is not null after performing join

## **Restoring Database**

The lab from now onwards will utilize the dvdrental database. Download the database tar file from <https://github.com/imkumaraju/dvdrental-sample-database>.

**CREATE DATABASE dvdrental;**

Use the pg\_restore tool to load data into the dvdrental database that we had just created as using the command:

You have to run this command outside the database

**pg\_restore -U postgres -d dvdrental [path\_to\_tar\_file]**

The next page contains the schema of dvdrental database.

